
Chapter 3: Arithmetic for Computers

Ngo Lam Trung



[with materials from *Computer Organization and Design, 4th Edition*,
Patterson & Hennessy, © 2008, MK
and M.J. Irwin's presentation, PSU 2008]

Content



- ❑ Integer representation and arithmetic
- ❑ Floating point number representation and arithmetic

Overview

- ❑ Computers store data as sequences of 1s and 0s
- ❑ How these sequences can be converted/displayed as audio, image, photo,...?



Type Name	Size
bool	8 bit unsigned
byte, unsigned char	8 bit unsigned
char	8 bit signed
unsigned int	16 bit unsigned
short, int	16 bit unsigned
unsigned long	32 bit unsigned
long	32 bit signed
float	32 bit floating point value
string	Array of byte



In this chapter: How to represent complicated data types in binary

Sign and unsigned integer

❑ Unsigned integer

$$\begin{aligned} \mathbf{x} &= \mathbf{x}_{n-1}\mathbf{x}_{n-2}\dots\mathbf{x}_1\mathbf{x}_0 \\ &= \mathbf{x}_{n-1}2^{n-1} + \mathbf{x}_{n-2}2^{n-2} + \dots + \mathbf{x}_12^1 + \mathbf{x}_02^0 \end{aligned}$$

- ❑ 32-bit unsigned numbers

0000	0000	0000	0000	0000	0000	0000	0000	0	$= 0_{\text{ten}}$
0000	0000	0000	0000	0000	0000	0000	0001	1	$= 1_{\text{ten}}$
...									
0111	1111	1111	1111	1111	1111	1111	1111	0	$= 2,147,483,646_{\text{ten}}$
0111	1111	1111	1111	1111	1111	1111	1111	1	$= 2,147,483,647_{\text{ten}}$
1000	0000	0000	0000	0000	0000	0000	0000	0	$= 2,147,483,648_{\text{ten}}$
1000	0000	0000	0000	0000	0000	0000	0001	1	$= 2,147,483,649_{\text{ten}}$
...									
1111	1111	1111	1111	1111	1111	1111	1111	0	$= 4,294,967,294_{\text{ten}}$
1111	1111	1111	1111	1111	1111	1111	1111	1	$= 4,294,967,295_{\text{ten}}$

MSB

LSB

Sign and unsigned integer

❑ Signed integer

$$\begin{aligned}
 X &= X_{n-1}X_{n-2}\dots X_1X_0 \\
 &= -X_{n-1}2^{n-1} + X_{n-2}2^{n-2} + \dots + X_12^1 + X_02^0
 \end{aligned}$$

❑ 32-bit signed numbers (2's complement):

	0000	0000	0000	0000	0000	0000	0000	0000	two	=	0	ten	
	0000	0000	0000	0000	0000	0000	0000	0001	two	=	+ 1	ten	
	...												
	0111	1111	1111	1111	1111	1111	1111	1110	two	=	+ 2,147,483,646	ten	
	0111	1111	1111	1111	1111	1111	1111	1111	two	=	+ 2,147,483,647	ten	<i>maxint</i>
	1000	0000	0000	0000	0000	0000	0000	0000	two	=	- 2,147,483,648	ten	
	1000	0000	0000	0000	0000	0000	0000	0001	two	=	- 2,147,483,647	ten	
	...												
MSB	1111	1111	1111	1111	1111	1111	1111	1110	two	=	- 2	ten	
	1111	1111	1111	1111	1111	1111	1111	1111	two	=	- 1	ten	<i>minint</i>

sign bit

LSB

Addition and subtraction

□ Addition

- Similar to what you do to add two numbers by hand
- Digits are added bit by bit from right to left
- Carries passed to the next digit to the left

□ Subtraction

- Negate the second operand then add to the first operand

$$\begin{array}{r} + \quad 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0111_{\text{two}} = 7_{\text{ten}} \\ \quad 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0110_{\text{two}} = 6_{\text{ten}} \\ \hline \quad 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 1101_{\text{two}} = 13_{\text{ten}} \end{array}$$

Examples

- ❑ All numbers are 8-bit signed integer

$$12 + 8 =$$

$$122 + 8 =$$

$$122 + 80 =$$

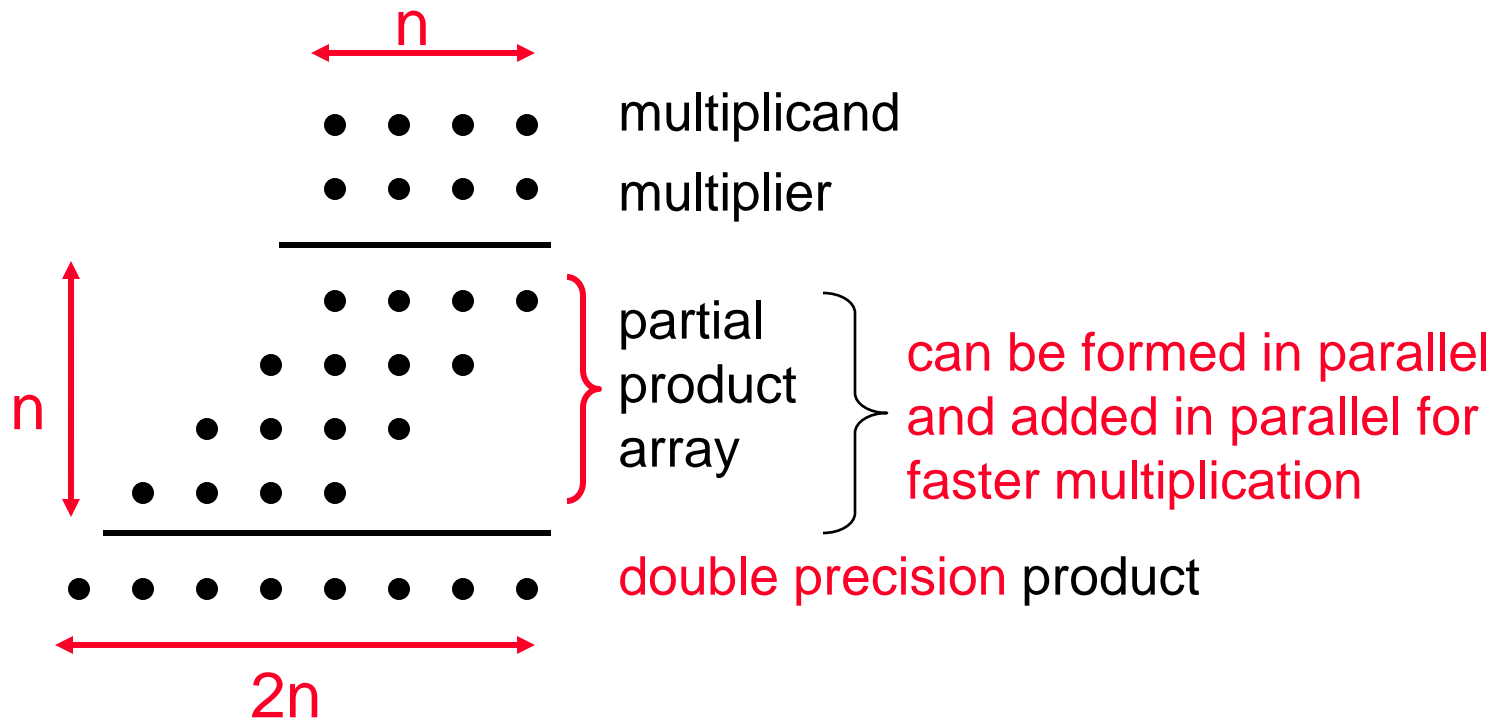
Dealing with Overflow

- ❑ Overflow occurs when the result of an operation cannot be represented in 32-bits, i.e., when the sign bit contains a **value** bit of the result and not the proper **sign** bit
 - ❑ When adding operands with different signs or when subtracting operands with the same sign, overflow can *never* occur

Operation	Operand A	Operand B	Result indicating overflow
A + B	≥ 0	≥ 0	< 0
A + B	< 0	< 0	≥ 0
A - B	≥ 0	< 0	< 0
A - B	< 0	≥ 0	≥ 0

Multiply

- ❑ Binary multiplication is just a *bunch* of right shifts and adds

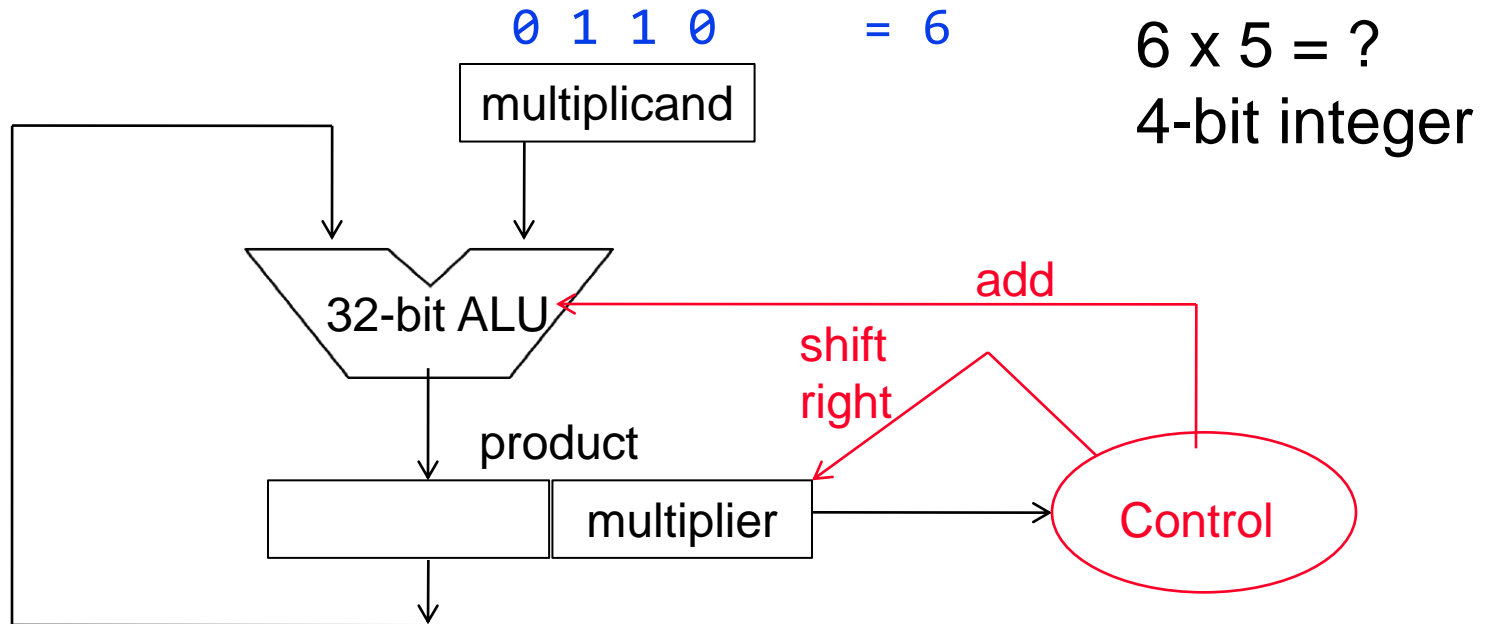


n -bit multiplicand and multiplier $\rightarrow 2n$ -bit product

Example

Multiplicand		1000 _{ten}
Multiplier	x	1001 _{ten}
		<hr/>
		1000
		0000
		0000
		1000
		<hr/>
Product		1001000 _{ten}

Add and Right Shift Multiplier Hardware



	0 0 0 0	0 1 0 1	= 5
add	0 1 1 0	0 1 0 1	LSB=1 → add multiplicand
	0 0 1 1	→ 0 0 1 0	shift right
add	0 0 1 1	0 0 1 0	LSB=0 → no change
	0 0 0 1	→ 1 0 0 1	shift right
add	0 1 1 1	1 0 0 1	LSB=1 → add multiplicand
	0 0 1 1	→ 1 1 0 0	shift right
add	0 0 1 1	1 1 0 0	LSB=0 → no change
	0 0 0 1	→ 1 1 1 0	shift right = 30

MIPS Multiply Instruction

- ❑ Multiply (`mult` and `multu`) produces a double precision product (2 x 32 bit)

```
mult    $s0, $s1        # hi || lo = $s0 * $s1
```


0	16	17	0	0	0x18
---	----	----	---	---	------

- Two additional registers: **hi** and **lo**
 - Low-order word of the product is left in processor register `lo` and the high-order word is left in register `hi`
 - Instructions `mfhi rd` and `mflo rd` are provided to move the product to (user accessible) registers in the register file
- ❑ Multiplies are usually done by fast, dedicated hardware and are much more complex (and slower) than adders

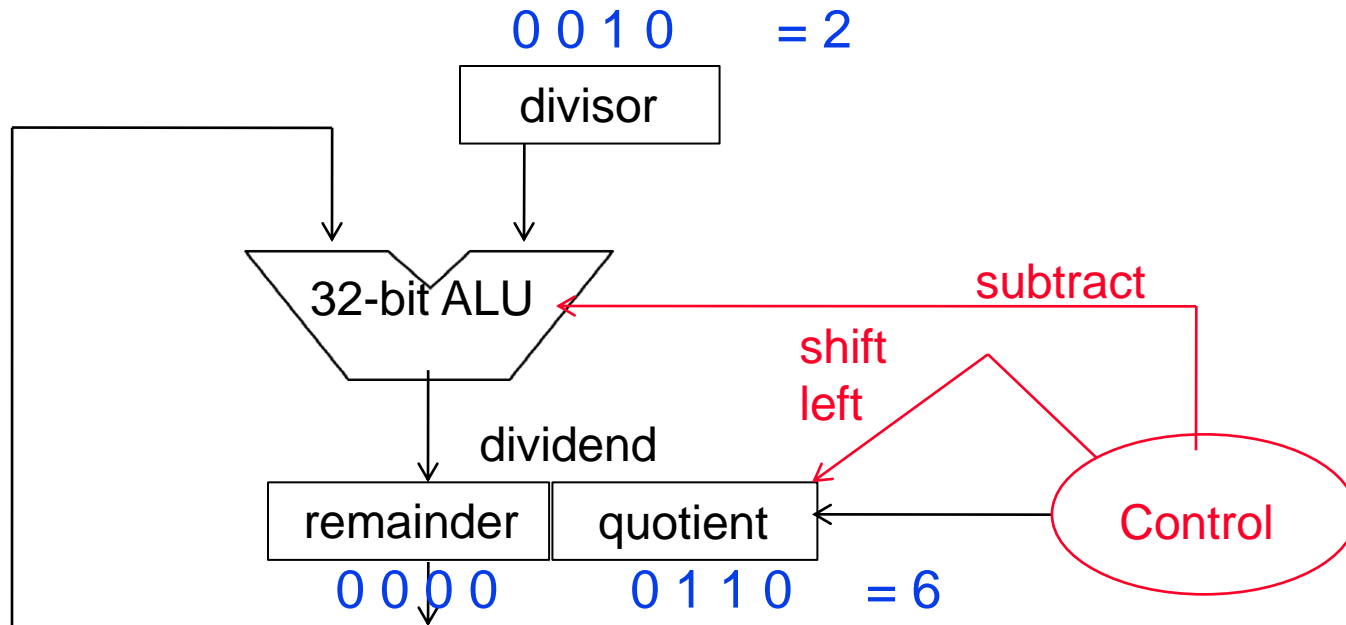
Division

- Division is just a *bunch* of quotient digit guesses and left shifts and subtracts

$$\text{dividend} = \text{quotient} \times \text{divisor} + \text{remainder}$$



Left Shift and Subtract Division Hardware



	0 0 0 0	←	1 1 0 0	
sub	1 1 1 0		1 1 0 0	rem neg, so 'ient bit = 0
	0 0 0 0		1 1 0 0	restore remainder
	0 0 0 1	←	1 0 0 0	
sub	1 1 1 1		1 1 0 0	rem neg, so 'ient bit = 0
	0 0 0 1		1 0 0 0	restore remainder
	0 0 1 1	←	0 0 0 0	
sub	0 0 0 1		0 0 0 1	rem pos, so 'ient bit = 1
	0 0 1 0	←	0 0 1 0	
sub	0 0 0 0		0 0 1 1	rem pos, so 'ient bit = 1
				= 3 with 0 remainder

MIPS Divide Instruction

- ❑ Divide (`div` and `divu`) generates the remainder in `hi` and the quotient in `lo`

`div $s0, $s1 # lo = $s0 / $s1`

`# hi = $s0 mod $s1`

0	16	17	0	0	0x1A
---	----	----	---	---	------

- Instructions `mfhi rd` and `mflo rd` are provided to move the quotient and remainder to (user accessible) registers in the register file
- ❑ As with multiply, divide ignores overflow so software must determine if the quotient is too large. Software must also check the divisor to avoid division by 0.

Representing Big (and Small) Numbers

- ❑ What if we want to encode the approx. age of the earth?

4,600,000,000 or 4.6×10^9

or the weight in kg of one a.m.u. (atomic mass unit)

[illegible]

or a famous number

PI = 3.14159....

There is no way we can encode either of the above in a 32-bit integer.

- ➔ We need reals or floating point numbers!

- ➔ Floating point numbers in decimal:

- 1000

- 1×10^3

- 0.1×10^4

Floating point number

- ❑ In decimal system

$$2013.1228 = 201.31228 * 10$$

$$= 20.131228 * 10^2$$

$$= 2.0131228 * 10^3$$

$$= 20131228 * 10^{-4}$$

- ❑ What is the “standard” form?

$$2.0131228 * 10^3 = \underline{2.0131228} \underline{E+03}$$

mantissa

exponent

- ❑ In binary $X = \pm 1.xxxxx * 2^{yyyy}$


- ❑ ***Sign, mantissa, and exponent need to be represented***

Floating point number

❑ Floating point representation in binary

$$(-1)^{\text{sign}} \times 1.F \times 2^{E-\text{bias}}$$

- Still have to fit everything in 32 bits (single precision)
- Bias = 127 with single precision floating point number



1 sign bit

8 bits

23 bits

❑ Defined by the IEEE 754-1985 standard

- ❑ Single precision: 32 bit
- ❑ Double precision: 64 bit
- ❑ Correspond to float and double in C

Examples

❑ Ex1: convert X into decimal value

$X = 1\mathbf{100\ 0001\ 0}101\ 0110\ 0000\ 0000\ 0000\ 0000$

sign = 1 \rightarrow X is negative

$E = 1000\ 0010 = 130$

$F = 10101100\dots00$

$$\begin{aligned}\rightarrow X &= (-1)^1 \times 1.101011\mathbf{000\dots00} \times 2^{130-127} \\ &= -1.101011 \times 2^3 = -1101.011 \\ &= -13.375\end{aligned}$$

Example

❑ Ex2: find decimal value of X

X = 0011 1111 1000 0000 0000 0000 0000 0000

sign = 0

e = 0111 1111 = 127

m = 000...0000 (23 bit 0)

$X = (-1)^0 \times 1.00...000 \times 2^{127-127} = 1.0$

Example

- Ex3: find binary representation of $X = 9.6875$ in IEEE 754 single precision

Converting X to plain binary

$$9_{10} = 1001_2$$

$$0.6875 \times 2 = 1.375 \quad \rightarrow \text{get bit } 1$$

$$0.375 \times 2 = 0.75 \quad \rightarrow \text{get bit } 0$$

$$0.75 \times 2 = 1.5 \quad \rightarrow \text{get bit } 1$$

$$0.5 \times 2 = 1.0 \quad \rightarrow \text{get bit } 1$$



$$\rightarrow 9.6875_{10} = 1001.1011_2$$

Example

- Ex3: find binary representation of $X = 9.6875$ in IEEE 754 single precision

$$X = 9.6875_{(10)} = 1001.1011_{(2)} = 1.0011011 \times 2^3$$

Then

$$S = 0$$

$$e = 127 + 3 = 130_{(10)} = 1000\ 0010_{(2)}$$

$$m = 001101100\dots00 \text{ (23 bit)}$$

Finally

$$X = 0100\ 0001\ 0001\ 1011\ 0000\ 0000\ 0000\ 0000$$

Examples

□ $1.0_2 \times 2^{-1} =$


□ $100.75_{10} =$

Some special values

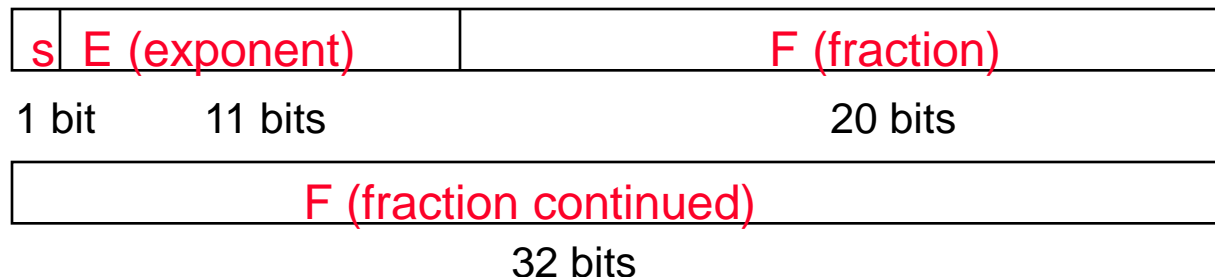
- ❑ Smallest+: 0 00000001 1.00000000000000000000000000000000
= $1 \times 2^{1-127}$
- ❑ Zero: 0 00000000 00000000000000000000000000000000
= true 0
- ❑ Largest+: 0 11111110 1.11111111111111111111111111111111
= $(2-2^{-23}) \times 2^{254-127}$

Too large or too small values

- ❑ **Overflow** (floating point) happens when a positive exponent becomes too large to fit in the exponent field
- ❑ **Underflow** (floating point) happens when a negative exponent becomes too large to fit in the exponent field



- ❑ One way to reduce the chance of underflow or overflow is to offer another format that has a larger exponent field
 - Double precision – takes two MIPS words



IEEE 754 FP Standard Encoding

- ❑ Special encodings are used to represent unusual events
 - \pm infinity for division by zero
 - NaN (not a number) for the results of invalid operations such as 0/0
 - True zero is the bit string all zero

Single Precision		Double Precision		Object Represented
E (8)	F (23)	E (11)	F (52)	
0000 0000	0	0000 ... 0000	0	true zero (0)
0000 0000	nonzero	0000 ... 0000	nonzero	\pm denormalized number
0111 1111 to +127,-126	anything	0111 ...1111 to +1023,-1022	anything	\pm floating point number
1111 1111	+ 0	1111 ... 1111	- 0	\pm infinity
1111 1111	nonzero	1111 ... 1111	nonzero	not a number (NaN)

Floating Point Addition

□ Addition (and subtraction)

$$(\pm F1 \times 2^{E1}) + (\pm F2 \times 2^{E2}) = \pm F3 \times 2^{E3}$$

- Step 0: Restore the hidden bit in F1 and in F2
- Step 1: **Align** fractions by right shifting F2 by $E1 - E2$ positions (assuming $E1 \geq E2$) keeping track of (three of) the bits shifted out in G R and S
- Step 2: **Add** the resulting F2 to F1 to form F3
- Step 3: **Normalize** F3 (so it is in the form 1.XXXXXX ...)
 - If F1 and F2 have the same sign $\rightarrow F3 \in [1,4) \rightarrow$ 1 bit right shift F3 and increment $E3$ (check for overflow)
 - If F1 and F2 have different signs \rightarrow F3 may require *many* left shifts each time decrementing $E3$ (check for underflow)
- Step 4: **Round** F3 and possibly **normalize** F3 again
- Step 5: Rehide the most significant bit of F3 before storing the result

Floating Point Addition Example

❑ Add: $0.5 + (-0.4375) = ?$

$$(0.5 = 1.0000 \times 2^{-1}) + (-0.4375 = -1.1100 \times 2^{-2})$$

- Step 0: Hidden bits restored in the representation above
- Step 1: Shift significand with the smaller exponent (1.1100) right until its exponent matches the larger exponent (so once)
- Step 2: Add significands
 $1.0000 + (-0.111) = 1.0000 - 0.111 = 0.001$
- Step 3: Normalize the sum, checking for exponent over/underflow
 $0.001 \times 2^{-1} = 0.010 \times 2^{-2} = \dots = 1.000 \times 2^{-4}$
- Step 4: The sum is already rounded, so we're done
- Step 5: Rehide the hidden bit before storing

Floating Point Multiplication

❑ Multiplication

$$(\pm F1 \times 2^{E1}) \times (\pm F2 \times 2^{E2}) = \pm F3 \times 2^{E3}$$

- Step 0: Restore the hidden bit in F1 and in F2
- Step 1: **Add** the two (biased) exponents and subtract the bias from the sum, so $E1 + E2 - 127 = E3$
also determine the sign of the product (which depends on the sign of the operands (most significant bits))
- Step 2: **Multiply** F1 by F2 to form a double precision F3
- Step 3: **Normalize** F3 (so it is in the form 1.XXXXXX ...)
 - Since F1 and F2 come in normalized $\rightarrow F3 \in [1,4) \rightarrow$ 1 bit right shift F3 and increment E3
 - Check for overflow/underflow
- Step 4: **Round** F3 and possibly **normalize** F3 again
- Step 5: Rehide the most significant bit of F3 before storing the result

Floating Point Multiplication Example

❑ Multiply

$$(0.5 = 1.0000 \times 2^{-1}) \times (-0.4375 = -1.1100 \times 2^{-2})$$

- Step 0: Hidden bits restored in the representation above
- Step 1: Add the exponents (not in bias would be $-1 + (-2) = -3$
and in bias would be $(-1+127) + (-2+127) - 127 = (-1-2) + (127+127-127) = -3 + 127 = 124$
- Step 2: Multiply the significands
 $1.0000 \times 1.110 = 1.110000$
- Step 3: Normalized the product, checking for exp over/underflow
 1.110000×2^{-3} is already normalized
- Step 4: The product is already rounded, so we're done
- Step 5: Rehide the hidden bit before storing

MIPS Arithmetic Logic Unit (ALU)

- ❑ Must support the Arithmetic/Logic operations of the ISA


add, addi, addiu, addu

sub, subu

mult, multu, div, divu

and, andi, nor, or, ori, xor, xori

beq, bne, slt, slti, sltiu, sltu



- ❑ With special handling for

- sign extend – addi, addiu, slti, sltiu
- zero extend – andi, ori, xori
- overflow detection – add, addi, sub